BrainAlpha: An Autonomous Multi-Agent System for Quantitative Alpha Discovery

Nuthdanai Wangpratham

QuantCorner Laboratory · February 27, 2026

Pioneering Innovation in Automated Investment Research

Abstract

Alpha decay is not a risk to be managed—it is an operational rate that must be outrun. Empirical estimates place the average decay rate of a quantitative signal in U.S. equities at 5.6% per year (McLean and Pontiff, 2016), implying that a portfolio of 1,000 live factors must replace 56–100 signals annually merely to remain static. No human research organization can sustain this throughput indefinitely. This paper presents **BrainAlpha**, an end-to-end autonomous multi-agent system that replicates the complete alpha discovery workflow of a professional quantitative researcher: hypothesis ideation, signal specification, expression generation, simulation-based evaluation, error recovery, and iterative repair—all without human intervention. The system comprises six specialized LangGraph modules (M0–M4, M3.5), coordinated through a shared Signal Alpha Object (SAO) that carries every candidate from raw investment thesis to Brain-validated alpha. The architecture introduces four innovations absent from prior systems: (i) a 660-cell structured exploration grid that transforms hypothesis search into an exhaustible coverage problem; (ii) a RAG-enhanced specification encoder that eliminates field hallucination at zero simulation cost; (iii) a failure-mode taxonomy (FM-1 through FM-6) that converts binary evaluation feedback into typed diagnostic signals; and (iv) a convergence-governed repair engine with five formal stop conditions that prevents wasted simulation budget on structurally unrecoverable expressions. All evaluation is grounded exclusively in WorldQuant Brain's production WebSim infrastructure.

Keywords: quantitative finance · alpha discovery · large language models · multi-agent systems · retrieval-augmented generation · WorldQuant Brain · automated research

1 Introduction

1.1 The Alpha Throughput Problem

Consider the following arithmetic. A systematic investment firm maintains a live portfolio of 1,000 uncorrelated alpha signals. The expected lifespan of each signal, given an empirically documented decay rate of 5.6% per year (McLean and Pontiff, 2016), is approximately 18 years. This sounds comfortable—until one accounts for post-publication arbitrage, factor crowding, and microstructure adaptation. In practice, signals with high academic visibility decay at double the baseline rate; those based on widely traded factors may lose statistical significance within five years. The net result: the firm must retire and replace 56–100 signals per year simply to hold its portfolio steady.

At a productivity rate of 2–5 viable ideas per researcher-week with a 15–30% quality screening pass rate (Tang et al., 2025), sustaining a 1,000-signal portfolio requires the continuous effort of dozens of dedicated quantitative researchers—a talent pool that is expensive, scarce, and increasingly competitive to recruit.

Central claim. *Alpha discovery is not an irreducibly human activity. It is a structured cognitive task that is decomposable into well-defined sub-tasks, each within the capability of current LLM technology—and amenable to automation at a scale and cost structure no human organization can match.*

1.2 Large Language Models as Quantitative Researchers

For three decades, automated alpha discovery relied on genetic programming—evolutionary search that scales computationally but lacks economic intuition. Without prior knowledge of which regions of the expression space contain meaningful signals, evolutionary methods spend the majority of their budget on formulae no human researcher would consider.

Large language models change this calculus. Pre-trained on decades of financial literature and academic asset pricing research, frontier LLMs have internalized the causal structure of financial markets to a measurable degree. A recent large-scale study evaluating 161,000 stock-day observations across the Russell 1000 found that an agentic LLM—with no curated data inputs, no fine-tuning, and no look-ahead bias—delivered **0.185% daily alpha** ($t = 2.51$) after Fama-French six-factor risk adjustment, corresponding to approximately **46% annualized excess return** at a Sharpe ratio of 2.43 (Chen, 2026).

System-level evidence reinforces this foundation (Table 1). Alpha-GPT (Wang et al., 2025) achieved a top-10 global ranking at WorldQuant's International Quant Championship 2024 among 41,000+ teams with a 284% IC improvement over unaided human researchers. AlphaAgent (Tang et al., 2025) achieved an 81% improvement in pass rate via AST-similarity regularization. QuantaAlpha (Han et al., 2026), current state of the art as of February 2026, achieves IC = 0.1501, ICIR = 0.9110, and 27.75% annualized return on CSI300.

1.3 WorldQuant Brain as Evaluation Oracle

A critical prerequisite for autonomous alpha research is a rigorous, standardized evaluation environment. In-house backtesting requires data infrastructure, survivorship bias controls, and simulation expertise—barriers that impede the rapid iteration LLM-based generation demands.

WorldQuant Brain removes this barrier. Brain provides the same simulation infrastructure underlying WorldQuant's live strategies, evaluating expressions against up to 3,000 stocks with a standardized battery of machine-readable quality checks: Sharpe ratio, fitness score, turnover bounds, drawdown constraints, and self-correlation limits. These checks are deterministic, interpretable, and constitute precisely the structured feedback that an autonomous optimization loop requires.

This creates an architectural separation of concerns central to BrainAlpha: the LLM subsystem is responsible exclusively for *creativity*; Brain is responsible exclusively for *truth*. The agent generates; Brain evaluates. Neither subsystem contaminates the other.

1.4 Contributions

C1. Structured exploration grid (M0): 660-cell Cartesian grid transforming hypothesis search into an exhaustible coverage problem with priority-weighted selection.

C2. RAG-enhanced specification (M1): grounds every hypothesis in Brain's live field catalogue, eliminating field hallucination at zero simulation cost.

C3. Failure-mode taxonomy (FM-1 through FM-6): converts binary simulation feedback into typed, actionable diagnostic signals for targeted repair.

C4. Convergence-governed repair engine (M4): five formal stop conditions prevent wasted budget on structurally unrecoverable expressions.

C5. Error recovery layer (M3.5): escalating three-strategy repair sequence intercepts API failures before repair budget is consumed.

2 Related Work

Table 1: Comparison of LLM-based alpha mining systems on four key capabilities. ✓ = addressed; ○ = partial; × = not addressed.

System	Key Contribution	Fully Autonomous	Structured Feedback	Field-Grounded	Production Evaluation
Alpha-GPT (Wang et al., 2025)	Human-AI interactive mining	×	×	○	○
AlphaAgent (Tang et al., 2025)	AST diversity regularization	✓	×	○	×
RD-Agent (Microsoft Research, 2025)	Full research loop agent	✓	×	×	×
QuantaAlpha (Han et al., 2026)	Trajectory-level evolution	✓	×	○	×
BrainAlpha (ours)	FM awa re-pai pip	✓	✓	✓	✓

The unresolved gap. No prior system addresses structured, failure-mode-aware iterative repair within a production evaluation environment. Prior systems treat evaluation as a black box returning pass or fail. BrainAlpha treats evaluation as a diagnostic instrument returning typed, actionable failure signals—and builds a repair engine architected around that signal structure.

3 System Architecture

3.1 Design Principles

P1 — Separation of Concerns.
LLM modules (M0, M1, M2, M4) handle *creativity*; the evaluation oracle (M3, Brain) handles *truth*. Neither subsystem contaminates the other.

P2 — Stateful Observability.
All pipeline state is persisted at every LangGraph node transition, enabling deterministic replay from any checkpoint without re-execution.

P3 — Structured Feedback over Binary Retry.
Failures are classified into typed failure modes before repair. Targeted strategies are selected from a structured catalogue indexed by failure class.

P4 — Data-Grounded Generation.
M1 retrieves verified Brain field names before expression generation. M2 enforces Brain's documented operator vocabulary via a deterministic validator. Field hallucination is eliminated before any simulation budget is consumed.

3.2 Pipeline Topology

BrainAlpha is a directed graph with conditional routing. Figure 1 illustrates the complete topology.

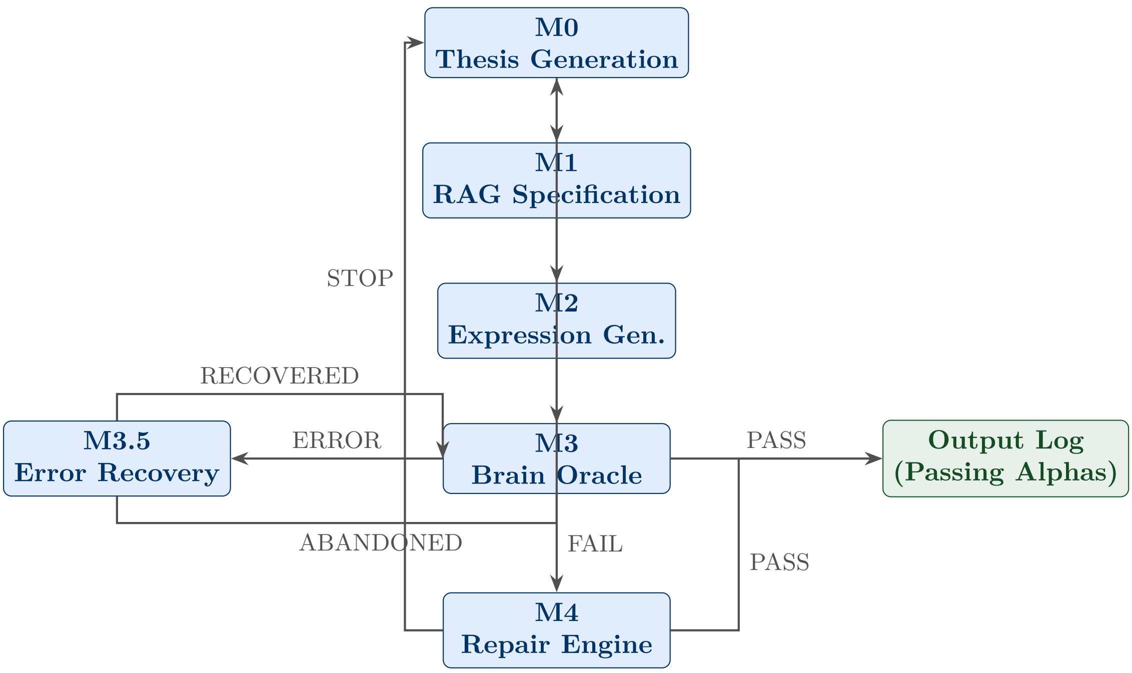

Figure 1: BrainAlpha pipeline topology. Every branch terminates: passing alphas accumulate in the output log; unrecoverable candidates return to M0 for fresh generation.

3.3 The Signal Alpha Object

All inter-module data flows through a single shared state object: the **Signal Alpha Object (SAO)**. The SAO is a typed dictionary that evolves as a candidate progresses through the pipeline, encoding the complete provenance of every candidate signal

(Table 2).

Table 2: SAO field groups, representative fields, and owning module.

Group	Key Fields	Owner	Stage
Identity	`id`, `source_agent`, `cell`	M0	Generation
Hypothesis	`statement`, `category`, `horizon`	M0	Generation
Specification	`primary_fields`, `neutralization`, `lookback_range`	M1	Specification
Expression	`expression`, `complexity`, `diversity_score`	M2	Generation
Simulation	`alpha_id`, `sharpe`, `fitness`, `checks[]`	M3	Evaluation
Evaluation	`status`, `failure_mode`, `m4_eligible`	M3	Evaluation
Repair	`repair_iterations`, `stop_reason`	M4	Repair

3.4 Failure Mode Taxonomy

Table 3: BrainAlpha Failure Mode Taxonomy (FM-1 through FM-6). Each FM code maps to a structured repair action catalogue in M4.

Code	Name	Diagnostic Signal	Primary Repair
FM-1	Low Signal Quality	Sharpe $<$ threshold, flat PnL	Neutralization, lookback extension
FM-2	Excessive Turnover	Turnover $>$ upper bound	Decay smoothing, lookback increase
FM-3	High Drawdown	Max drawdown $>$ bound	Winsorization, volatility conditioning
FM-4	Signal Reversal	Negative annualized returns	Negation (`-1 *`), inversion
FM-5	Low Coverage	Long/short count below minimum	`pasteurize()`, filter relaxation
FM-6	Low Turnover	Turnover $<$ lower bound	Lookback reduction, cross-sectional variation

4 Methodology

4.1 M0: Priority-Weighted Structured Exploration

4.1.1 The Exploration Grid

The hypothesis space is discretized into a **660-cell exploration grid**:

$$\text{Grid} = \underbrace{11}_{\text{Signal Family}} \times \underbrace{5}_{\text{Operator Pattern}} \times \underbrace{3}_{\text{Horizon}} \times \underbrace{4}_{\text{Conditioning}} = 660 \text{ cells}$$

Each cell maintains a lifecycle status from $\{$*empty*, *assigned*, *explored*, *exhausted*$\}$, enforcing breadth-first coverage. The consequence: alpha discovery is no longer a search over an infinite space. It is an exhaustible coverage problem over 660 well-defined cells.

4.1.2 Cell Priority Scoring

$$\text{Priority}(c, a) = 0.5 \cdot \text{Novelty}(c) + 0.3 \cdot \text{ExpectedYield}(c) + 0.2 \cdot \text{AgentAlignment}(c, a) \quad (1)$$

$\text{Novelty}(c) \in \{1.00, 0.50, 0.25, 0.00\}$ decreases monotonically with exploration status. $\text{ExpectedYield}(c)$ encodes a family-level prior on pass rates updated online after each confirmed passing alpha (+0.10 per pass in the same cell). $\text{AgentAlignment}(c, a) \in$

$\{0.3, 0.5, 0.7, 1.0\}$ captures semantic fit between cell taxonomy and the requesting agent's domain expertise.

4.2 M1: RAG-Enhanced Specification

The translation from thesis to formal specification is where naive LLM generation most frequently fails: models generate expressions referencing plausible-sounding but non-existent data fields. This *field hallucination* problem is structurally unavoidable—Brain's field catalogue evolves continuously, and no model can remain current through pretraining alone.

M1 embeds Brain's live field catalogue ($\approx$20,000 verified field names) into a vector index and retrieves the top-k most relevant fields for each incoming thesis. Retrieved fields are injected into the specification prompt as the exclusive permissible data vocabulary.

The output is then validated by a six-rule deterministic checker enforcing: non-empty primary fields; lookback range validity ($1 \leq \min \leq \max \leq 120$); recognized neutralization strategy; non-empty operator family; non-empty direction; hypothesis length (≥ 20 characters). Validation is zero-cost: no LLM invocation, no simulation budget consumed.

4.3 M2: Diversity-Enforced Expression Generation

M2 generates candidate expressions using Chain-of-Thought prompting at temperature $T = 0.7$ and applies a Jaccard diversity gate. Expression e_{new} is accepted only if:

$$\max_{e \in \mathcal{A}} \frac{|\text{ops}(e_{\text{new}}) \cap \text{ops}(e)|}{|\text{ops}(e_{\text{new}}) \cup \text{ops}(e)|} \; < \; 0.70 \tag{2}$$

where $\text{ops}(e)$ is the multi-set of operator tokens in expression e and $\mathcal{A}$ is the set of accepted expressions. This enforces structural diversity—not merely lexical variation.

4.4 M4: Convergence-Governed Repair

4.4.1 Phase 1: Recoverability Assessment

Before any repair iteration, three hard-stop conditions trigger immediate termination *without LLM invocation*:

I1. SELF_CORRELATION = FAIL: structural duplication of a live signal. Expression-level repairs cannot resolve idea-level duplication.

I2. Reversed complex signal: `returns` ≤ 0 and `complexity` > 20 AST nodes. Negation cannot propagate through a deeply nested expression without operator-semantics conflicts.

I3. Excessive failures: $n_{\text{failing}} \geq 4$. No single-action repair can address four independent deficiencies.

A soft stop applies when: $\max_m \text{gap\_pct}(m) > 0.50$, where $\text{gap\_pct}(m) = |v_m - \tau_m| \, / \, |\tau_m|$.

4.4.2 Phase 2: ReAct Repair Agent

For candidates cleared by the recoverability gate, a **ReAct** (Reasoning + Acting) agent autonomously diagnoses the failure, selects a repair strategy, and rewrites the expression within a single multi-step reasoning loop. The agent is equipped with three tools:

- `diagnose_checks` — ranks all failing Brain checks by gap severity (largest percentage gap from threshold first).
- `lookup_strategies` — retrieves untried repair actions from the structured catalogue for a given check name.
- `draft_expression` — invokes an LLM sub-agent to apply one targeted repair instruction and return the revised expression.

The agent reasons through the problem—*What is most broken? What has already been tried? What is the least disruptive fix?*—before issuing tool calls. At each iteration, the agent has access to the full session repair history: every expression attempted, every check result, and every action already applied are visible in its context, preventing redundant repairs. Practitioners may also inject **priority rules** via `M4Config.custom_instructions`—domain knowledge such as "always negate the signal first for FM-4 reversals"—which the agent treats as higher-priority guidance before consulting the catalogue.

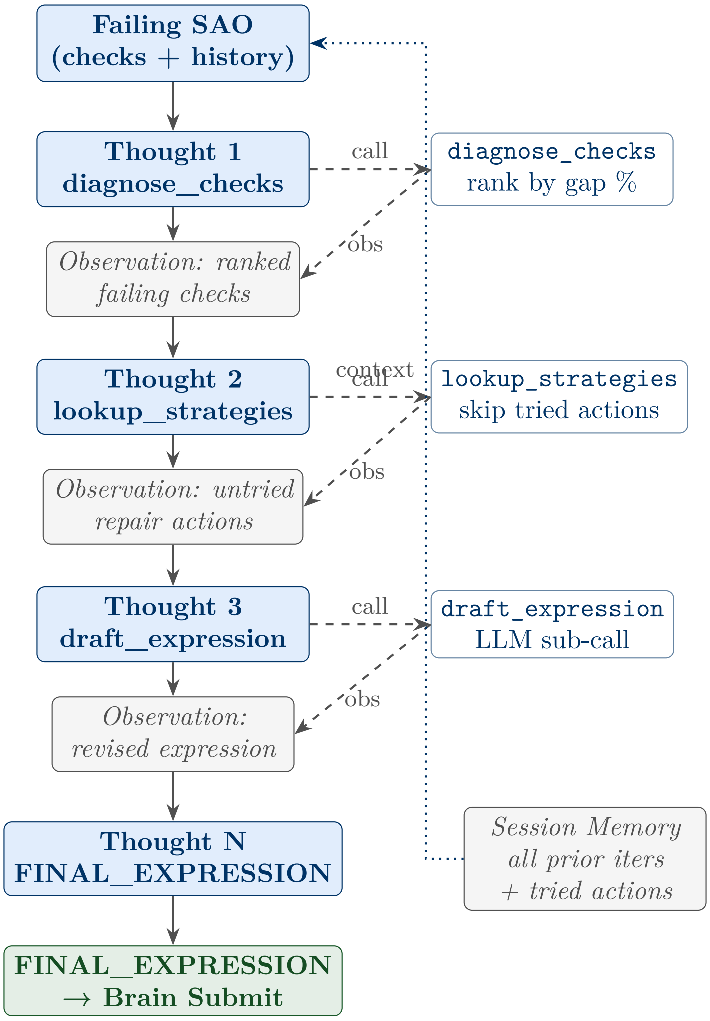

Figure 2: ReAct Agent internal loop (M4 Phase 2). Each "Thought" step reasons about the current state before issuing a tool call. Session memory (all prior iterations, expressions, and attempted actions) is injected as context at every step, preventing redundant repairs. The loop terminates when the agent emits `FINAL_EXPRESSION`, which is then submitted to Brain WebSim.

4.4.3 Convergence Governance

The loop terminates when any of five conditions is met: (i) the expression passes all checks; (ii) maximum iteration count reached; (iii) per-iteration Sharpe improvement $< \delta = 0.02$ for three consecutive iterations; (iv) gap percentage grows rather than shrinks (divergence); or (v) all catalogue actions exhausted.

5 Discussion

5.1 Structure Matters More Than Scale

The central methodological lesson of BrainAlpha is that *structure matters more than scale* in LLM-based generation tasks with formal correctness requirements. The exploration grid, the RAG-grounded specification encoder, the deterministic validator, the failure-mode classifier, and the ReAct repair agent each independently reduce failure rates; together, they produce a system whose output quality is substantially higher than naive LLM prompting—even with the same underlying model.

This principle generalizes: any domain where LLM generation must satisfy formal constraints benefits from the same pattern: constrain the generation space first, then use the LLM to explore within that constrained space with domain-informed intelligence.

5.2 The Economics of Alpha Industrialization

At the throughputs demonstrated in initial deployment cycles, a BrainAlpha instance running continuously against Brain's WebSim API can generate validated alpha signals at a rate rivaling a small human research team—at a marginal cost approaching zero per additional signal.

The implication is significant. Alpha discovery is not an irreducibly human activity. The question for any quantitative organization is not *whether* to automate alpha discovery, but *when* the cost of building the automation infrastructure will be recovered by the throughput gains it delivers. Given current performance estimates and the documented cost of quantitative research talent, the crossover point is substantially shorter than most practitioners assume.

5.3 Human-AI Complementarity

BrainAlpha does not propose to replace human quantitative researchers. The system reflects an explicit division of labor: the LLM operates within a structured space defined by human domain expertise; human researchers define the space. Researchers can focus on activities requiring genuine judgment—evaluating economic coherence, assessing regime sensitivity, managing portfolio-level risk—while BrainAlpha handles high-volume expression generation, evaluation, and repair at throughputs no human can match.

6 Conclusion

This paper has presented BrainAlpha—an autonomous multi-agent pipeline for quantitative alpha discovery operating end-to-end against WorldQuant Brain's production infrastructure.

The system's four principal innovations—the structured exploration grid, the RAG-enhanced specification encoder, the failure-mode taxonomy, and the convergence-governed repair engine—together constitute a system that is not merely technically interesting but economically consequential. Alpha decay is a quantifiable throughput deficit. BrainAlpha is a quantifiable throughput generator. The gap between the two is the business case for automated alpha discovery—and it is a gap that only grows larger as systematic investment management scales.

"Alpha decay waits for no one. Neither should the research pipeline."

References

Zefeng Chen (2026). Agentic AI as autonomous equity analyst: Generating actionable investment signals at scale. *Under review*.

Fama, E. F. (1970). Efficient capital markets: A review of theory and empirical work. *Journal of Finance*, 25(2):383–417.

Han, Y. et al. (2026). QuantaAlpha: Trajectory-level evolutionary optimization for quantitative alpha mining. *arXiv preprint*.

Harvey, C. R., Liu, Y., and Zhu, H. (2016). . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1):5–68.

Jegadeesh, N. and Titman, S. (1993). Returns to buying winners and selling losers. *Journal of Finance*, 48(1):65–91.

Kakushadze, Z. (2016). 101 formulaic alphas. *Wilmott*, 2016(84):72–80.

McLean, R. D. and Pontiff, J. (2016). Does publishing research destroy stock return predictability? *Journal of Finance*, 71(1):5–32.

Microsoft Research (2025). RD-Agent: Towards data-centric automatic R&D. *NeurIPS 2025*.

Tang, Y. et al. (2025). AlphaAgent: Anti-decay alpha mining with LLM agents. *KDD 2025*.

Wang, X. et al. (2025). Alpha-GPT 2.0: Human-AI interactive alpha mining. *AAAI 2025*.